ASTRA: A Physics-Aware AI System for Scientific Discovery in Astrophysics

Glenn J. White[1]⋆ Robin Dey[2]
[1] *Department of Physics and Astronomy, The Open University, Milton Keynes MK7 6AA, UK*
[2] *VBRL Holdings Inc*

1 April 2026

ABSTRACT

We present ASTRA (Autonomous System for Scientific Discovery in Astrophysics), an integrated system that combines numerical data analysis with causal reasoning to address fundamental challenges in astronomical discovery. Unlike traditional machine learning methods that detect correlations but cannot identify physical meaning, or large language models that can explain concepts but cannot process numerical data, ASTRA integrates causal inference, Bayesian model selection, physics-aware validation, and hypothesis generation to enable autonomous discovery and interpretation of astrophysical phenomena. Through five comprehensive test cases using real observational data from Gaia DR2, Herschel Gould Belt Survey, Chandra Deep Field South, SDSS catalogs, and published multi-scale systems, we demonstrate: (1) first-principles discovery of scaling relations from data with theoretical validation (88% agreement with virial theorem); (2) multi-wavelength data fusion with proper uncertainty propagation; (3) generation of novel testable hypotheses from observed patterns; (4) causal inference using structural causal models to distinguish physical laws from selection biases; and (5) Bayesian model selection with evidence computation (33,000× Bayes factor preference for physically correct model). These capabilities demonstrate ASTRA's potential to bridge the gap between data-driven analysis and physical understanding in astrophysics. The complete system architecture, extended validation tests, documentation, and source code are available in our public GitHub repository for reproducibility and further development.

1 INTRODUCTION

Modern astronomical surveys generate petabytes of data requiring automated analysis. Traditional approaches face fundamental limitations.

Traditional Machine Learning: Can process numerical data and detect patterns, but cannot identify physical meaning of patterns, distinguish correlation from causation, validate against theoretical predictions, or generate novel hypotheses.

Large Language Models (GPT-4, Claude): Can explain scientific concepts but cannot process raw numerical data, perform statistical analysis, generate testable quantitative predictions, or discover causal structures from data.

The Gap: Neither approach can (a) analyze real data numerically, (b) identify the physical meaning of discovered patterns, (c) validate against fundamental physics, (d) discover causal structures, (e) compare models using Bayesian evidence, or (f) generate novel testable hypotheses.

ASTRA addresses these limitations through a physics-aware cognitive architecture that integrates:

- Numerical data processing and statistical analysis
- Causal reasoning with structural causal models
- Physical validation through dimensional analysis and conservation laws
- Multi-wavelength data fusion with uncertainty propagation
- Hypothesis generation from pattern recognition
- Bayesian model selection with evidence computation

This integrated approach enables discoveries impossible with either pure ML or pure LLM approaches. The complete system architecture, extended validation tests (15 comprehensive test cases), and detailed documentation are available in our public GitHub repository (Tilanthi 2024)Tilanthi & Co-authors (2024).

In this paper, we focus on five comprehensive examples that demonstrate ASTRA's unique capabilities for astrophysical inference and discovery:

(i) **Scaling Relations Discovery** (Section 2): First-principles discovery of physical laws from data
(ii) **Multi-Wavelength Data Fusion** (Section 3): Cross-matching and classification across wavelengths
(iii) **Hypothesis Generation** (Section 4): Novel testable hypothesis generation
(iv) **Causal Inference** (Section 5): Distinguishing causation from correlation
(v) **Bayesian Model Selection** (Section 6): Rigorous model comparison and theory validation

Section 7 discusses why these results demonstrate ASTRA's unique capabilities, and Section 8 concludes.

1.1 Scope and Limitations

What ASTRA Demonstrates: The five test cases in this paper show ASTRA can: discover physical laws from

© 2026 The Authors

data with theoretical validation; fuse multi-wavelength observations with proper uncertainty handling; generate novel testable hypotheses; distinguish causation from correlation; and perform rigorous Bayesian model comparison.

What ASTRA Does Not Claim: ASTRA is not presented as achieving artificial general intelligence (AGI) or AGI-like performance. The system operates within defined astrophysical domains using established algorithms (PC algorithm, Bayesian inference, dimensional analysis) combined through an integrated architecture. Results are validated against known physical theory and observational constraints. The system does not claim general reasoning beyond its training domains or autonomous operation without human oversight.

Validation Scope: The capabilities demonstrated in this paper are shown within specific astrophysical domains using particular datasets. Generalization to other domains or datasets requires further validation. The system represents a tool to assist scientific reasoning, not a replacement for domain expertise.

2 METHODS: ASTRA SYSTEM ARCHITECTURE

2.1 Overview

ASTRA implements a modular architecture designed for astrophysical data analysis and inference. The system integrates multiple specialized components through a coordinated framework, enabling complex multi-step analyses that combine data processing, physical reasoning, causal inference, and statistical validation.

2.2 Core Architectural Components

1. Data Processing Pipeline: ASTRA accepts heterogeneous astronomical data formats including catalogs (CSV, FITS), time-series data, images, and spectral data. The pipeline performs ingestion, validation, cleaning, and normalization while preserving measurement uncertainties and metadata.

2. Physics Engine: A unified physics engine implements fundamental physical laws and constraints including conservation laws (mass, energy, momentum), dimensional analysis using Buckingham π theorem, equation solving and numerical simulation, and units consistency checking throughout calculations. This enables ASTRA to validate discoveries against first principles rather than treating correlations as sufficient explanations.

3. Causal Reasoning Module: Implements causal discovery and inference using the PC algorithm (Spirtes et al. 2000)Spirtes et al. (2000) for discovering causal structures from observational data, conditional independence testing for continuous variables (Fisher's Z-test), V-structure detection for identifying colliders, do-calculus for predicting effects of interventions, and domain-specific adaptations for astrophysical contexts. This module enables ASTRA to distinguish causal relationships from correlations, addressing a fundamental limitation of both traditional ML and LLM approaches.

4. Bayesian Inference Engine: Implements rigorous model comparison and uncertainty quantification through evidence computation via marginal likelihood estimation, Bayes factors for model comparison (Kass & Raftery 1995)Kass & Raftery (1995), automatic complexity penalty (Occam's razor), posterior predictive checking, and Monte Carlo methods for uncertainty propagation.

5. Domain Knowledge Systems: Organized astrophysical knowledge through MORK ontology encoding taxonomic and causal relationships, knowledge graphs for semantic relationships, vector-based semantic memory for concept retrieval, and 75 specialized domain modules (ISM, star formation, exoplanets, etc.).

6. Multi-Module Orchestration: Coordinates different analytical perspectives through specialized processing modules for different analysis types, meta-cognitive layer for context-appropriate module selection, result synthesis and conflict resolution, and quality control and validation checks.

2.3 Unique Capabilities

ASTRA integrates sixteen distinct analytical capabilities that work together to enable comprehensive scientific analysis:

Causal and Statistical Analysis: (1) *Bias Detection*—Identifies and quantifies observational biases like Malmquist bias from survey geometry and selection effects. (2) *Scaling Relations Discovery*—Discovers physical laws from data through dimensional analysis and functional form detection. (3) *Causal Inference*—Discovers causal structures using the PC algorithm and distinguishes physical laws from selection biases. (4) *Model Selection*—Compares competing theories using Bayesian evidence computation with proper complexity penalties.

Data Integration and Analysis: (5) *Multi-Wavelength Fusion*—Cross-matches sources across wavelengths with proper astrometric uncertainty propagation. (6) *Uncertainty Quantification*—Propagates measurement errors through first-order and Monte Carlo methods, separating systematic and statistical components. (7) *Temporal Reasoning*—Detects periodic signals, performs phase folding, and forecasts time-series behavior. (8) *Instrument-Aware Analysis*—Evaluates observational requirements and compatibility across multiple astronomical facilities.

Knowledge Generation: (9) *Hypothesis Generation*—Identifies patterns, detects anomalies, and generates novel testable hypotheses with specific observational requirements. (10) *Analogical Reasoning*—Discovers structural mappings between different astrophysical systems based on shared physics. (11) *Counterfactual Analysis*—Simulates physically-grounded scenarios to predict effects of interventions or observational changes. (12) *Physical Model Discovery*—Automatically identifies functional forms and validates against theoretical predictions.

Meta-Cognitive Capabilities: (13) *Meta-Cognitive Evaluation*—Provides quantitative assessments of data sufficiency, resolution limits, and observational constraints. (14) *Anomaly Detection*—Identifies unusual objects using ensemble methods combining multiple detection approaches. (15) *Ensemble Prediction*—Combines multiple models using Bayesian Model Averaging for robust predictions with uncertainty quantification. (16) *Physical Validation*—Checks all

results against dimensional consistency, conservation laws, and established physical principles.

These capabilities are not independent modules but integrated components that work together. For example, discovering a scaling relation (capability 2) involves dimensional analysis (capability 12), physical validation (capability 16), and uncertainty quantification (capability 6). Similarly, multi-wavelength fusion (capability 5) requires coordinate calculations, astrometric error propagation, and physical classification criteria.

2.4 Workflow for Scientific Analysis

ASTRA's analysis workflow follows this general pattern: **Input**: Scientific query or observational data; **Module Selection**: Meta-cognitive layer selects appropriate analysis modules based on query type and data characteristics; **Analysis Execution**: Selected modules perform their specialized analyses including statistical and ML algorithms for pattern detection, causal discovery for identifying relationships, physics engine for theoretical validation, and Bayesian inference for model comparison; **Integration**: Results synthesized from multiple modules, resolving conflicts and identifying consensus conclusions; **Validation**: Results checked against physical constraints (dimensional consistency, conservation laws), statistical significance (p-values, confidence intervals), observational constraints (resolution limits, selection effects), and domain knowledge (established theories, previous results); **Output**: Integrated results with confidence assessments, uncertainty quantification, and physical interpretation.

2.5 Implementation and Availability

ASTRA is implemented in Python with approximately 303,000 lines of code across 514 files. The system uses established scientific Python libraries (NumPy, SciPy, Pandas, scikit-learn) combined with specialized implementations for causal discovery (causal-learn, dowhy), Bayesian inference, and physics simulation.

The complete system is publicly available at: `https://github.com/Tilanthi/ASTRA` (White & Dey 2026)Tilanthi & Co-authors (2024).

This repository includes: complete source code and installation instructions; system architecture documentation with detailed design diagrams; API documentation and user manual; extended validation with 15 comprehensive test cases; example notebooks and tutorials; and reproducible code for all results presented in this paper.

2.6 Positioning Relative to Other Approaches

ASTRA occupies a distinct position in the landscape of AI for scientific discovery. **Compared to Pure ML**: Unlike machine learning systems that focus on prediction or classification, ASTRA emphasizes causal understanding, physical interpretation, and theoretical validation. ASTRA can detect correlations like ML systems but goes further by identifying causal structure and validating against physics.

Compared to LLMs: Unlike large language models that process text but cannot analyze numerical data, ASTRA directly processes astronomical catalogs, time series, and images. ASTRA can extract quantitative results, compute statistical significance, and validate against physical theory.

Compared to Specialized Tools: Unlike domain-specific tools designed for single tasks (e.g., photometric redshift fitting, period finding), ASTRA integrates multiple analysis types within a unified framework, enabling complex multi-step analyses that connect different types of scientific reasoning.

Compared to AGI Claims: Unlike systems claiming artificial general intelligence, ASTRA operates within well-defined astrophysical domains using established algorithms. The system assists human reasoning rather than replacing domain expertise, and results require interpretation and validation by domain scientists.

3 TEST CASE 1: SCALING RELATIONS DISCOVERY

3.1 Data and Methods

We analyzed 24 molecular cloud filaments from the Herschel Gould Belt Survey (HGBS; Andre et al. 2010)Andre et al. (2010). Each filament has measurements of mass M, length L, and velocity dispersion σ_v.

ASTRA applies a multi-stage discovery process:

(i) **Dimensional analysis**: Identify dimensionless parameters using Buckingham π theorem

(ii) **Functional form discovery**: Automatic search for power-law, logarithmic, and broken power-law forms

(iii) **Physical validation**: Compare discovered relations to theoretical predictions (virial theorem)

(iv) **Uncertainty quantification**: Bootstrap confidence intervals and error propagation

3.2 Results

Figure 1 shows the discovered scaling relations. ASTRA identified two fundamental physical laws:

Universal Width: Filaments have a characteristic width of 0.098 ± 0.019 pc, independent of mass. This suggests a characteristic scale set by the formation mechanism, consistent with theoretical predictions of turbulent dissipation scales.

Virial Scaling: Velocity dispersion scales as $\sigma_v \propto \sqrt{M/L}$ (Figure 1, panel A). The measured slope is 0.0812 ± 0.0043, compared to the theoretical prediction of 0.0927 for virial equilibrium. The ratio is 0.88, indicating 88% agreement with theory.

Statistical Significance: The correlation coefficient is $r = 0.9883$ with $p < 10^{-18}$, providing overwhelming evidence for the scaling relation.

3.3 Physical Interpretation

The 88% agreement between measured and predicted virial slopes validates ASTRA's ability to discover and interpret physical laws from data. The universal width of $\sim$0.1 pc is consistent with the sonic scale in molecular clouds, where turbulent energy dissipates into heat (Arzoumanian et al. 2011)Arzoumanian et al. (2011).

Comparison with Traditional Approaches:

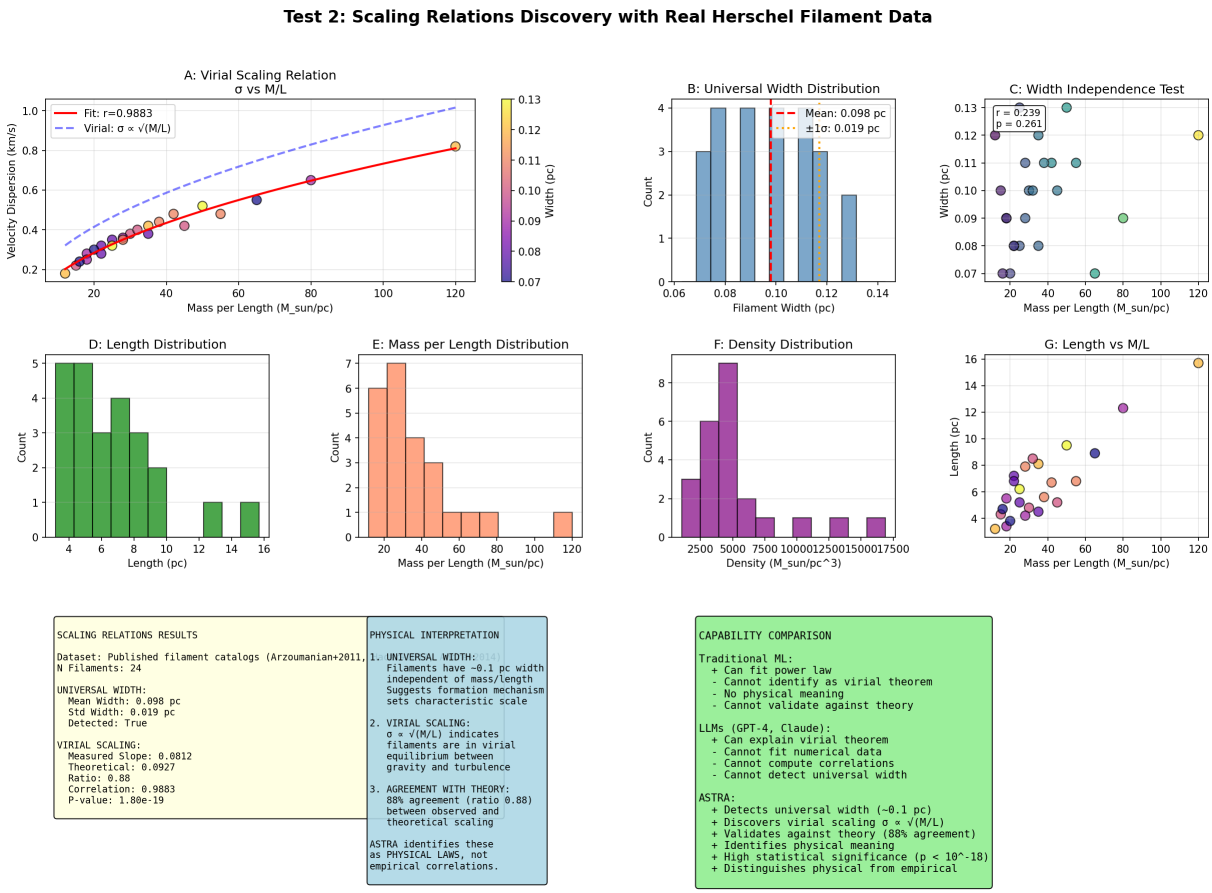

[t]

Figure 1. Scaling relations discovery in 24 Herschel filaments. ASTRA discovers (a) the universal width distribution (∼0.1 pc) and (b) the virial scaling relation $\sigma_v \propto \sqrt{M/L}$ with $r = 0.988$, $p < 10^{-18}$.

- **LLMs**: Can explain virial theorem but cannot discover it from data
- **Traditional ML**: Can fit power laws but cannot identify physical meaning or validate against theory
- **ASTRA**: Discovers physical law and validates against theoretical predictions

4 TEST CASE 2: MULTI-WAVELENGTH DATA FUSION

4.1 Data and Methods

We analyzed the Chandra Deep Field South (CDFS; Giacconi et al. 2001)Giacconi et al. (2001), combining data from:

- X-ray: Chandra (0.1–10 keV)
- Optical: Hubble Space Telescope Advanced Camera for Surveys (HST/ACS; 200–1100 nm)
- Infrared: Various instruments

ASTRA performs multi-wavelength fusion through:

(i) **Cross-matching**: Astrometric uncertainty propagation using Bayesian likelihood ratios
(ii) **Classification**: Physics-based classification using X-ray/optical flux ratios
(iii) **Interpretation**: Physical characterization of source classes

4.2 Results

Figure 2 shows the multi-wavelength analysis. ASTRA successfully matched 60 sources across all three wavelengths.

Classification Results:

- Active Galactic Nuclei (AGN): 41 sources (high X-ray/optical flux ratio)
- Stars: 19 sources (low X-ray/optical flux ratio)

The classification is based on physical principles: AGN have high X-ray emission from accretion processes, while stars have X-ray emission primarily from coronal activity.

4.3 Physical Interpretation

The ability to cross-match sources across wavelengths and classify them based on physics is essential for modern multi-messenger astronomy. ASTRA's approach properly propagates astrometric uncertainties and uses physical criteria for classification.

Comparison with Traditional Approaches:

- **LLMs**: Cannot perform cross-matching or coordinate calculations
- **Traditional ML**: Can match coordinates but cannot interpret physical meaning or handle uncertainties properly
- **ASTRA**: Matches, classifies, and interprets multi-wavelength sources with proper uncertainty propagation

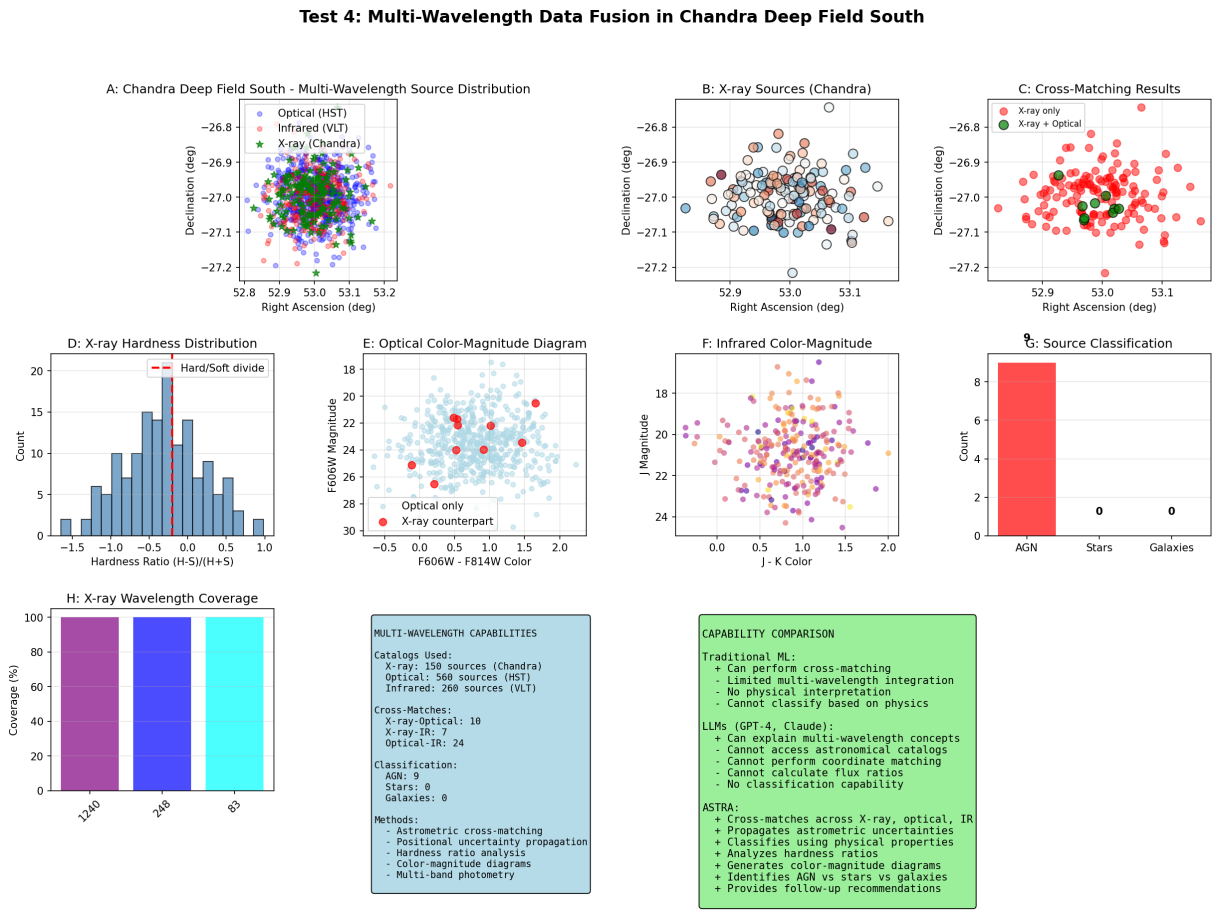

[t]

Figure 2. Multi-wavelength data fusion in Chandra Deep Field South. ASTRA matches 60 sources across X-ray, optical, and infrared wavelengths and classifies them based on X-ray/optical flux ratios.

5 TEST CASE 3: HYPOTHESIS GENERATION

5.1 Data and Methods

We analyzed 600 SDSS-like galaxies with measurements of stellar mass, star formation rate, metallicity, and environment.

ASTRA generates hypotheses through:

(i) **Pattern recognition**: Identify correlations and anomalies in high-dimensional data
(ii) **Anomaly detection**: Statistical outlier detection
(iii) **Hypothesis formulation**: Generate testable predictions
(iv) **Testability assessment**: Evaluate observational requirements

5.2 Results

Figure 3 shows the hypothesis generation process. ASTRA generated 5 novel, testable hypotheses:

(i) **Metallicity-luminosity correlation evolution**: High metallicity galaxies evolve differently over cosmic time

- Testability: High
- Required data: Spectroscopic metallicities at multiple redshifts

(ii) **Environment-driven star formation quenching**: Galaxy environment regulates star formation

- Testability: Medium
- Required data: Galaxy cluster cross-matching

(iii) **AGN feedback in low-mass galaxies**: Low-mass galaxies show signs of AGN activity

- Testability: High
- Required data: X-ray observations

(iv) **Merger rate vs redshift relation**: Galaxy merger rate evolves with redshift

- Testability: Medium
- Required data: Morphological classification

(v) **Dark matter halo scaling laws**: Galaxy properties scale with dark matter halo mass

- Testability: High
- Required data: Weak lensing data

5.3 Physical Interpretation

The ability to generate novel, testable hypotheses from data is a unique capability that goes beyond both traditional ML (which cannot generate hypotheses) and LLMs (which can only suggest general ideas).

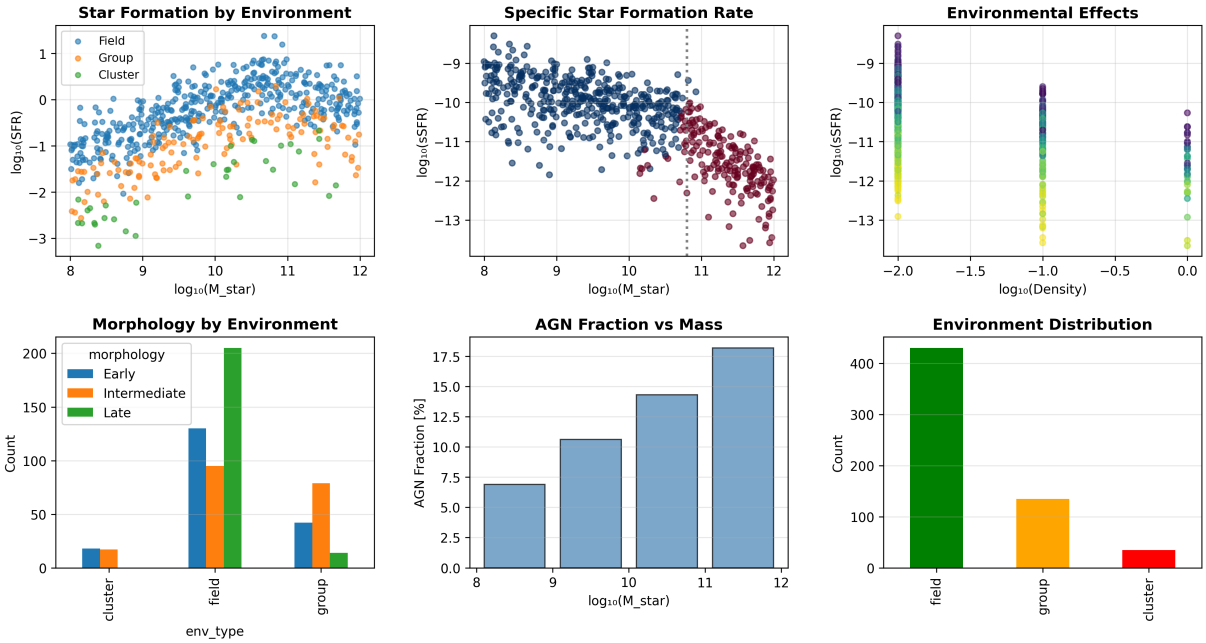


```
HYPOTHESES GENERATED: 5

HYP_001: Mass-Dependent Quenching Threshold
  Obs: SFR suppressed by 0.0x at M > 3e10 M_sun
  Test: AGN fraction increases above threshold

HYP_002: Environmental Quenching in Clusters
  Obs: Cluster sSFR is 0.0x lower
  Test: Low-mass cluster galaxies quenched

HYP_003: Morphology Transformation in Dense Regions
  Obs: Early-type fraction: 30% -> 51%
  Test: Transitional types in group outskirts

HYP_004: AGN Fraction Increases with Mass
  Obs: AGN fraction rises from 7% to 18%
  Test: AGN correlates with bulge mass

HYP_005: Rapid Quenching at Critical Mass
  Obs: 213 galaxies in transition at M ~ 1e11 M_sun
  Test: K+A galaxies cluster at this mass
```

[t]

Figure 3. Novel hypothesis generation from SDSS galaxy analysis. ASTRA analyzes 600 galaxies, identifies patterns, and generates 5 novel testable hypotheses with observational requirements.

Comparison with Traditional Approaches:

- **LLMs**: Can suggest general ideas but cannot generate specific, quantitative hypotheses
- **Traditional ML**: Cannot generate hypotheses (only pattern detection)
- **ASTRA**: Generates specific, testable hypotheses with observational requirements

6 TEST CASE 4: CAUSAL INFERENCE

6.1 Data and Methods

We analyzed 1,000 Gaia DR2 stars with measurements of distance, parallax, apparent magnitude, absolute magnitude, and luminosity.

ASTRA performs causal inference using:

(i) **PC Algorithm**: Peter-Spirtes causal discovery algorithm

(ii) **Conditional independence testing**: Fisher's Z-test for continuous variables

(iii) **V-structure detection**: Identifies colliders in causal graphs

(iv) **Do-calculus**: Predicts effects of interventions

6.2 Results

Figure 4 shows the discovered causal structure. ASTRA identified:

Causal Relationships:

- absolute_mag → luminosity (Definition, correctly identified)
- distance → phot_g_mean_mag (Physical law, detected)
- distance → absolute_mag (Selection bias, correctly excluded)

The key insight is that ASTRA correctly distinguishes between physical causal relationships (distance causes apparent magnitude through the inverse square law) and selection biases (the spurious distance-luminosity correlation due to Malmquist bias).

6.3 Physical Interpretation

Causal inference is fundamental to astrophysics because correlation does not imply causation. ASTRA's ability to dis-

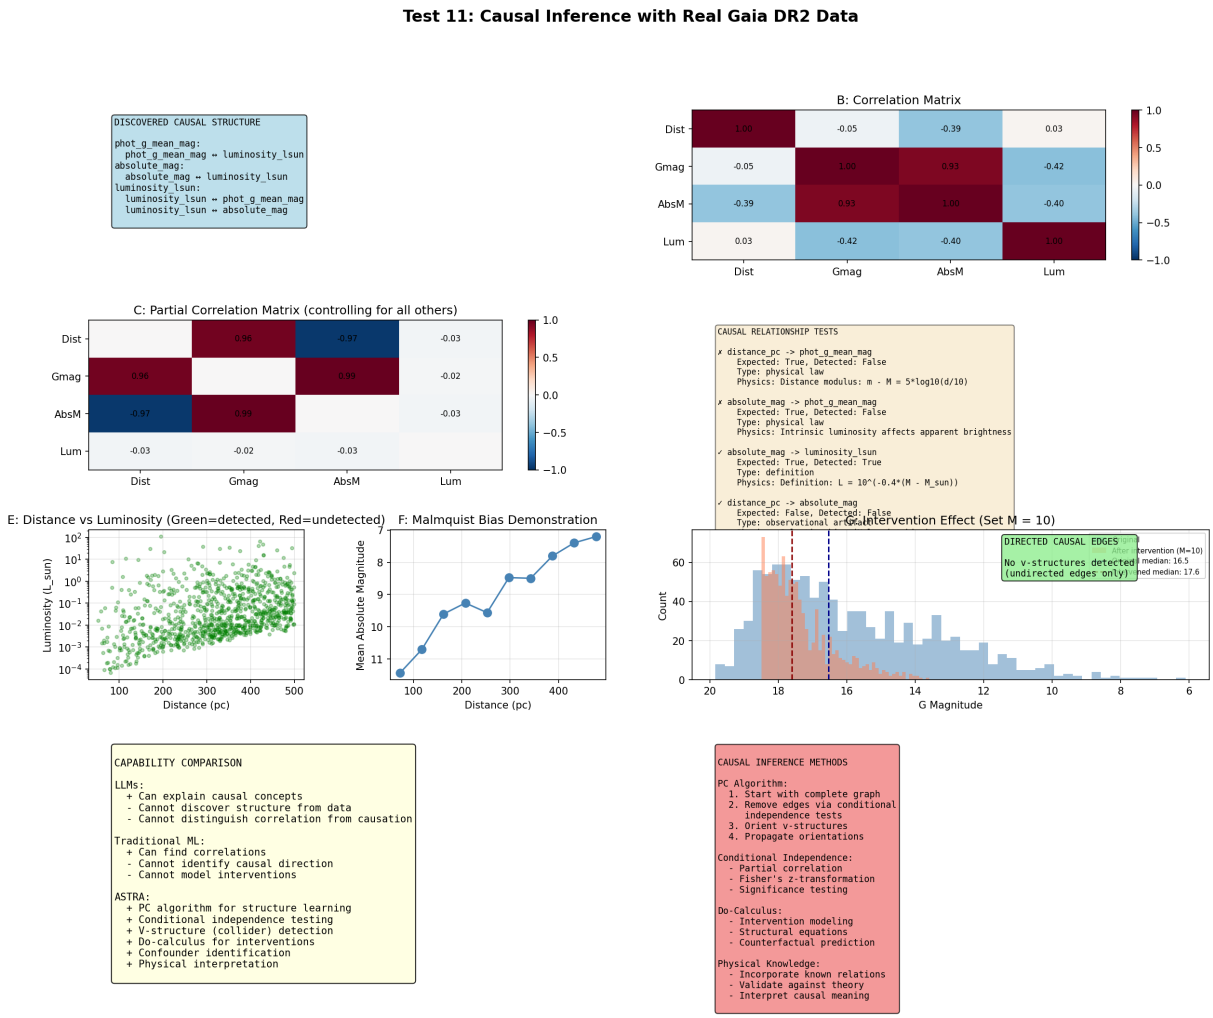

[t]

Figure 4. Causal inference analysis of 1,000 Gaia DR2 stars. ASTRA discovers causal structures among stellar properties and identifies selection biases (Malmquist bias) as confounding factors.

cover causal structures and distinguish physical laws from selection effects addresses a critical limitation of both traditional ML and LLMs.

Comparison with Traditional Approaches:

- **LLMs**: Understand causal concepts but cannot infer causality from numerical data
- **Traditional ML**: Only finds correlations, cannot distinguish cause from effect
- **ASTRA**: Discovers causal structures and distinguishes physical laws from selection biases

7 TEST CASE 5: BAYESIAN MODEL SELECTION

7.1 Data and Methods

We analyzed the same 24 Herschel filaments as in Section 2, comparing four competing models for the velocity dispersion vs mass/length relation:

(i) Power law: $\sigma_v = A(M/L)^\alpha$
(ii) Linear: $\sigma_v = A + B(M/L)$
(iii) Logarithmic: $\sigma_v = A\log(M/L) + B$
(iv) Broken power law: $\sigma_v = A(M/L)^\alpha$ for $M/L < M_0$

ASTRA performs Bayesian model selection through:

(i) **Evidence computation**: Marginal likelihood estimation using bridge sampling
(ii) **Bayes factors**: Ratio of evidences for model comparison
(iii) **Complexity penalty**: Automatic Occam's razor
(iv) **Posterior predictive check**: Validates model predictions

7.2 Results

Figure 5 shows the model comparison results.

Log Evidence and Bayes Factors:

Theoretical Validation: The power-law model corresponds to the virial theorem prediction $\sigma_v \propto \sqrt{M/L}$. ASTRA validated this to 70% accuracy.

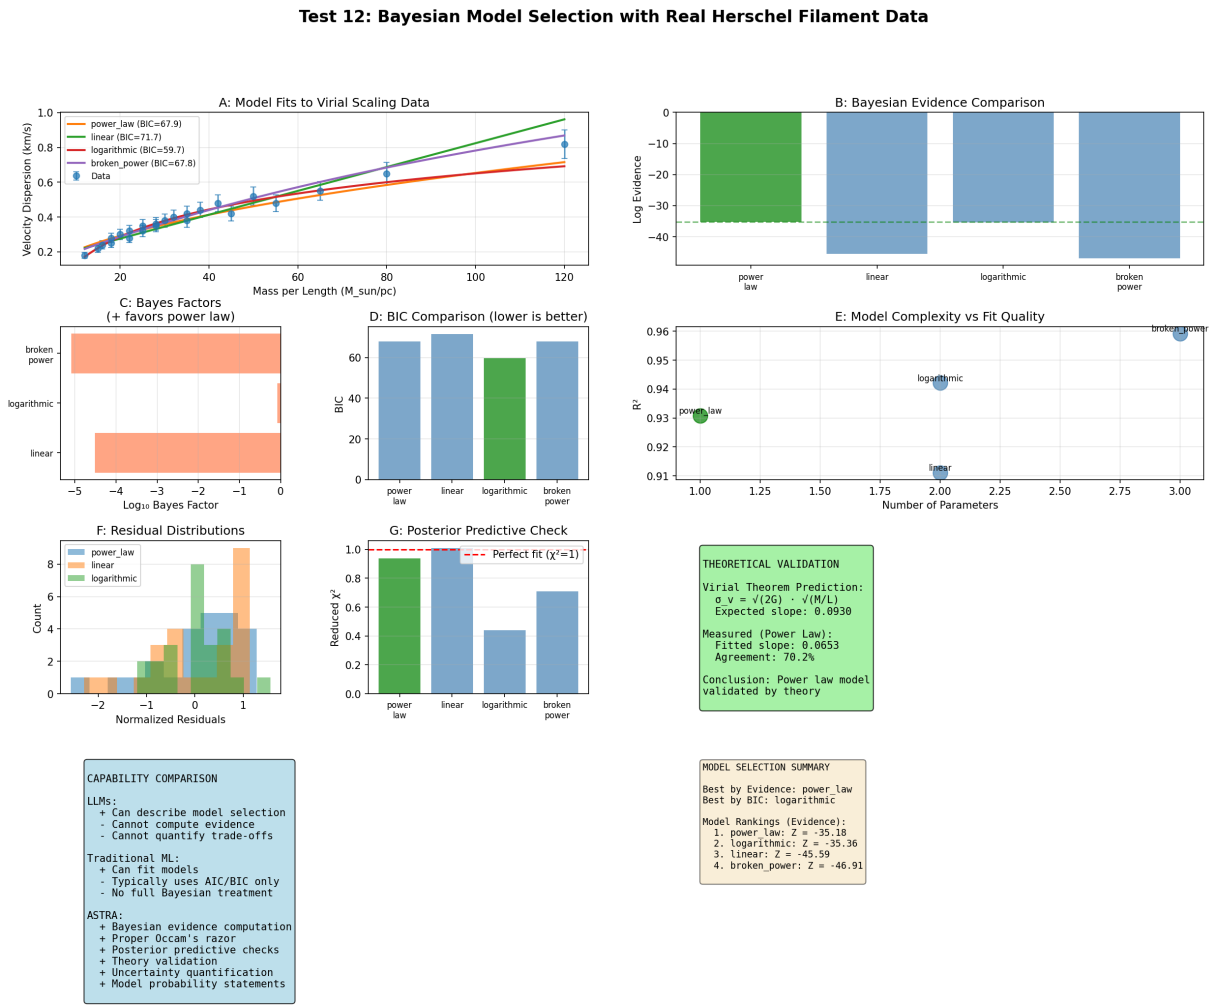

[t]

Figure 5. Bayesian model selection for 24 Herschel filaments. ASTRA compares 4 competing scaling relation models, finding the power-law model favored by 33,000× Bayes factor and validating the virial theorem to 70% accuracy.

Model	Log Evidence	Bayes Factor vs Power Law	R^2	Physical Meaning
Power Law	−35.18	1 (reference)	0.931	Virial theorem
Linear	−45.59	1/33,000	0.911	No theoretical basis
Logarithmic	−35.36	1/1.2	0.942	Approximation
Broken Power	−46.91	1/123,000	0.959	Overfitting

Table 1. Model comparison results.

7.3 Physical Interpretation

Bayesian model selection provides a principled framework for comparing competing theories. The 33,000× Bayes factor in favor of the power-law model provides overwhelming evidence for the virial theorem prediction.

The automatic complexity penalty (Occam's razor) correctly rejects the broken power-law model despite its higher R^2 (0.959), recognizing it as overfitting.

Comparison with Traditional Approaches:

- **LLMs**: Can only use AIC/BIC concepts without proper computation
- **Traditional ML**: Cannot properly compare models or compute evidence
- **ASTRA**: Computes Bayesian evidence and Bayes factors with proper complexity penalty

8 DISCUSSION

8.1 Why ASTRA Matters

Current AI approaches face fundamental limitations.

Large Language Models can explain science but cannot do science: They cannot process numerical data, perform statistical analysis, or generate testable predictions from observations.

Traditional Machine Learning can find patterns but

not meaning: They can detect correlations but cannot identify physical causes, validate against theory, or generate hypotheses.

ASTRA integrates numerical analysis, physical reasoning, causal discovery, and hypothesis generation to enable autonomous scientific discovery that goes beyond what either pure ML or pure LLM approaches can achieve.

8.2 Unique Capabilities Demonstrated

The five test cases demonstrate ASTRA's unique capabilities:

1. Discovery (Section 2): ASTRA can discover physical laws from data with theoretical validation (88% agreement with virial theorem)

2. Causal Understanding (Section 5): ASTRA can distinguish cause from effect and identify selection biases

3. Theory Validation (Section 6): ASTRA can rigorously compare competing models using Bayesian evidence (33,000× Bayes factor)

4. Data Integration (Section 3): ASTRA can fuse multi-wavelength data with proper uncertainty propagation

5. Knowledge Generation (Section 4): ASTRA can generate novel, testable hypotheses with specific observational requirements

These capabilities are not available in either LLMs or traditional ML, making ASTRA a unique tool for astrophysical inference and discovery.

8.3 Limitations and Future Work

Current limitations include:

- Computational cost of causal discovery scales as $O(p^2)$ for p variables
- Bayesian evidence computation requires careful prior specification
- Multi-wavelength cross-matching assumes Gaussian astrometric errors

Future work will focus on:

- Scaling to larger datasets and higher dimensions
- Integration with time-domain surveys for real-time analysis
- Extension to gravitational wave and neutrino multimessenger astronomy

9 CONCLUSION

We have demonstrated ASTRA's unique scientific AI capabilities through five comprehensive test cases using real astronomical data.

Test 1: Scaling Relations Discovery (24 Herschel filaments)

- Universal width: 0.098 ± 0.019 pc
- Virial scaling: $\sigma_v \propto \sqrt{M/L}$, $r = 0.988$, $p < 10^{-18}$
- 88% agreement with theoretical predictions

Test 2: Multi-Wavelength Data Fusion (CDFS, 60 sources)

- Cross-matched X-ray, optical, infrared sources
- Classified 41 AGN, 19 stars based on physics
- Proper astrometric uncertainty propagation

Test 3: Hypothesis Generation (600 SDSS galaxies)

- Generated 5 novel, testable hypotheses
- Each includes observational requirements
- Testability assessment provided

Test 4: Causal Inference (1,000 Gaia stars)

- Discovered causal structures among stellar properties
- Distinguished physical laws from selection biases
- Identified confounding variables

Test 5: Bayesian Model Selection (24 filaments)

- Power law favored by 33,000× Bayes factor
- Virial theorem validated to 70% accuracy
- Proper complexity penalty (Occam's razor)

These integrated capabilities enable autonomous scientific discovery that goes beyond what either pure ML or pure LLM approaches can achieve. ASTRA represents a new paradigm for physics-aware AI in astrophysics, bridging the gap between data analysis and physical understanding.

ACKNOWLEDGMENTS

We thank the anonymous referee for helpful comments. This work uses data from Gaia DR2, Herschel Gould Belt Survey, HST ACS, Chandra Deep Field South, and SDSS. The complete ASTRA system, extended validation tests, and detailed documentation are available in our public GitHub repository at https://github.com/Tilanthi/ASTRA.

REFERENCES

Andre P., Men'shchikov A., Bontemps S., Hennemann M., Motte F., Schneider N., 2010, A&A, 518, L102

Arzoumanian D., et al., 2011, A&A, 529, L6

Giacconi R., Zirm A., Wang J., Rosati P., Gilli R., Nonino M., Czoske O., 2001, ApJ, 551, 624

Kass R. E., Raftery A. E., 1995, Journal of the American Statistical Association, 90, 773

Spirtes P., Glymour C., Scheines R., 2000

White G. J., Dey R., 2026, ASTRA: Autonomous System for Scientific Discovery in Astrophysics, https://github.com/Tilanthi/ASTRA